

Language Modeling

Probabilistic Language Models

Probabilistic Language Models learn the probability of the next token, based on previous tokens.

Given

that a set of words $w_1, w_2, w_3, \dots, w_k$ predict the next word as w_{k+1}

$P(w_{k+1} | w_k, w_{k-1}, w_{k-2}, \dots)$ This is multi-class classification problem inspired by naive-bayes

N-gram models

An N-gram language model predicts the probability of a given N-gram within any sequence of words in the language.

Consider “I have a dream” exists within a corpus.

$P(\text{I have a dream}) = P(\text{I}) \cdot p(\text{have}) \cdot p(\text{a}) \cdot p(\text{dream})$

Eg: a tri-gram model:

< start >	I	want	a	dream	job	have	dog	company
(< start >, I)	0	0	0.0714	0	0	0	0.142	0
(I,want)	0	0	0	0.0714	0	0	0	0
(want,a)	0	0	0	0	0.0714	0	0	0
(a,dream)	0	0	0	0	0	0.0714	0	0
(dream,job)	0	0	0	0	0	0	0	0
(i,have)	0	0	0	0.142	0	0	0	0
(have,a)	0	0	0	0	0.0714	0.0714	0	0
(a,dog)	0	0	0	0	0	0	0	0
(dream,company)	0	0	0	0	0	0	0	0
(company, < / start>)	0	0	0	0	0	0	0	0
(job, < / start>)	0	0	0	0	0	0	0	0
(dog, < / start>)	0	0	0	0	0	0	0	0

Building N-grams: Maximum Likelihood Estimation

Maximum likelihood estimation maximizes a likelihood function, that, under the assumed statistical model, the observed data is most probable.

For n-grams, the likelihood function for the most probable word at position k is

$$P(w_k | w_{k-1}, w_{k-2} \dots w_1) = \frac{\text{Count}(w_k, w_{k-1})}{\text{Count}(w_{k-1}, w_{k-2} \dots w_1)}$$

N-gram models are trained by building a dictionary that given a set of n-1 words, predicts nth word by maximum likelihood estimation. Thus, n-gram assumes that

N-gram assumption: The probability of each word is independent, given the words before it, and only depends on the fraction of their occurrence within the corpus

Evaluation

To evaluate the performance of n-gram (or) autoregressive models, the following metrics are used

- **Average Loglikelihood:** of a text document from a corpus is

- $\log P_{eval}(text)_{avg} = \frac{1}{|word|} \sum_{word} \log P_{train}(word)$

- where $|word|$ is number of words in the text

- **Cross entropy:** of a text document from a corpus is negative of average log-likelihood

- $\text{Cross Entropy}(text) = -\log P_{eval}(text)_{avg}$

- **Perplexity:** of a text document from a corpus is defined as exponential of cross-entropy

- $\text{Perplexity}(text) = e^{\text{cross_entropy}(text)}$

- A lower perplexity indicates a higher model's understanding of a language

Accounting OOV (Out of vocabulary words)

To account for OOV words, the following algorithms are used

- Laplace smoothening helps to tackle problem of zero probability. We add an additional [UNK] token to training set and modify likelihood function

- **K-Smoothering:**

$$P(w_k | w_{k-1}, w_{k-2} \dots w_1) = \frac{\text{Count}(w_k, w_{k-1}) + k}{\text{Count}(w_{k-1}, w_{k-2} \dots w_1) + k \cdot V}$$

- where k is a randomly chosen hyperparameter and V is vocabulary size (Number of unique n-grams in the training corpus)

- **Add 1 Smoothening:**

$$P(w_k | w_{k-1}, w_{k-2} \dots w_1) = \frac{\text{Count}(w_k, w_{k-1}) + 1}{\text{Count}(w_{k-1}, w_{k-2} \dots w_1) + V}$$

- where V is vocabulary size (Number of unique n-grams in the training corpus). Add 1 smoothening is a simplified version of K-Smoothering where $K = 1$

- Effect of Smoothening

$$P_{train}([UNK]) = \frac{k}{V}$$

- Back-off and Interpolation

- **Back-off:** this technique falls back to lower-order N-gram models when a higher-order N-gram model is unavailable

$$P_{backoff}(w_n | w_{n-1}, w_{n-2} \dots w_1) = \{P_{train}(w_n | w_{n-1}, w_{n-2} \dots w_1) \text{Count}(w_n, w_{n-1} \dots w_1) \geq 0\} P_{backoff}(w_{n-1} | w_{n-2}, w_{n-3} \dots w_1)$$

- **Interpolation:** This technique combines the likelihoods of multiple N-gram models to estimate the next word. In linear interpolation, we typically take a weighted sum of likelihoods as a predictor for the next token

$$P(w_n | w_{n-1}, w_{n-2} \dots w_1) = \lambda_1 P(w_1) + \lambda_2 P(w_2 | w_1) + \dots \lambda_{n-1} P(w_{n-1} | w_{n-2}, w_{n-3} \dots w_1) + \lambda_n P(w_n | w_{n-1}, w_{n-2} \dots w_1)$$
